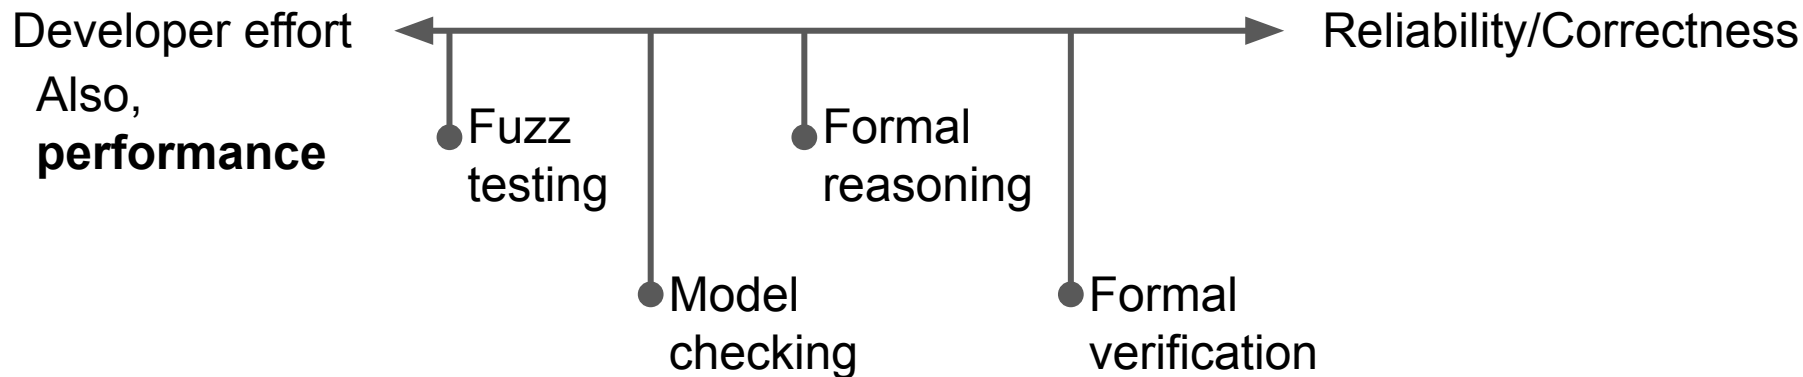


# Primer on formal verification: Bugs



# Primer on formal verification: Curry-Howard isomorphism

Direct relationship between programs and proofs:

Proposition  $\sim$  Type

Proof of a proposition  $\sim$  Program with a type signature

Coq: Interactive theorem prover.

- Model complex behavior. Example: Reading and writing to a disk.
- State propositions about the model's behavior. Example:  $A \Rightarrow B$ .
- Prove the proposition. Example: Construct a program with the type signature  $A \rightarrow B$ .
- Coq checks the proof by type-checking the program.


# Using Crash Hoare Logic for Certifying the FSCQ File System

Haogang Chen, Daniel Ziegler, Tej Chajed,  
Adam Chlipala, Frans Kaashoek, and Nickolai Zeldovich

MIT CSAIL

# File systems are complex and have bugs


File systems are complex (e.g., Linux ext4 is  $\sim 60,000$  lines of code) and have many bugs:



Cumulative number of patches for file-system bugs in Linux; data from [Lu et al., FAST'13]

# File systems are complex and have bugs

File systems are complex (e.g., Linux ext4 is  $\sim 60,000$  lines of code) and have many bugs:




Cumulative number of patches for file-system bugs in Linux; data from [Lu et al., FAST'13]

New file systems (and bugs) are introduced over time

# File systems are complex and have bugs

File systems are complex (e.g., Linux ext4 is ~60,000 lines of code) and have many bugs:



Cumulative number of patches for file-system bugs in Linux; data from [Lu et al., FAST'13]

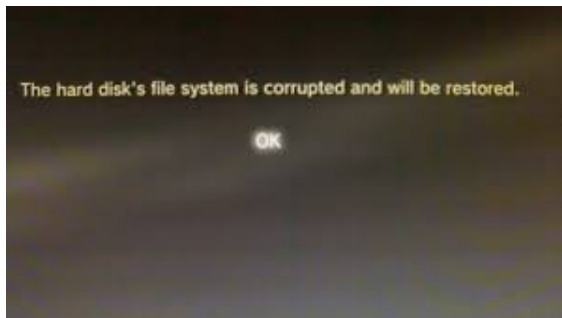
New file systems (and bugs) are introduced over time

Some bugs are serious: **security exploits**, **data loss**, etc.

# Reasoning about crashes is important

File system must recover from crash with data intact


- Crash due to power failure, hardware failures, or software bugs



## Risk: crashes expose subtle bugs

Crashes can expose that a file update involves *several* disk writes

Appending to a file requires three disk writes:






## Risk: crashes expose subtle bugs

Crashes can expose that a file update involves *several* disk writes


Appending to a file requires three disk writes:



## Risk: crashes expose subtle bugs

Crashes can expose that a file update involves *several* disk writes


Appending to a file requires three disk writes:



## Risk: crashes expose subtle bugs

Crashes can expose that a file update involves *several* disk writes


Appending to a file requires three disk writes:



## Risk: crashes expose subtle bugs

Crashes can expose that a file update involves *several* disk writes


Appending to a file requires three disk writes:



## Risk: crashes expose subtle bugs

Crashes can expose that a file update involves *several* disk writes

Appending to a file requires three disk writes:



# Developers struggle with partially-updated states

```
commit 353b67d8ced4dc53281c88150ad295e24bc4b4c5
--- a/fs/jbd/checkpoint.c
+++ b/fs/jbd/checkpoint.c
@@ -504,7 +503,25 @@ int cleanup_journal_tail(journal_t *journal)
     spin_unlock(&journal->j_state_lock);
     return 1;
 }
+ spin_unlock(&journal->j_state_lock);
+
+ /*
+  * We need to make sure that any blocks that were recently written out
+  * --- perhaps by log_do_checkpoint() --- are flushed out before we
+  * drop the transactions from the journal. It's unlikely this will be
+  * necessary, especially with an appropriately sized journal, but we
+  * need this to guarantee correctness. Fortunately
+  * cleanup_journal_tail() doesn't get called all that often.
+  */
+ if (journal->j_flags & JFS_BARRIER)
+     blkdev_issue_flush(journal->j_fs_dev, GFP_KERNEL, NULL);
+
+ spin_lock(&journal->j_state_lock);
+ if (!tid_gt(first_tid, journal->j_tail_sequence)) {
+     spin_unlock(&journal->j_state_lock);
+     /* Someone else cleaned up journal so return 0 */
+     return 0;
+ }
+ /* OK, update the superblock to recover the freed space.
+  * Physical blocks come first: have we wrapped beyond the end of
+  * the log? */
```

Mistakes cause data loss [Yang et al. 2006, Pilai et al. 2014, Zheng et al. 2014]

## Goal: certify a complete file system under crashes

- A file system with a **machine-checkable proof**
- that its implementation meets **its specification**
- under **normal execution**
- and under any sequence of **crashes**
- including **crashes during recovery**



# Contributions

**CHL:** Crash Hoare Logic for persistent storage

- Crash condition and recovery semantics
- CHL automates parts of proof effort
- Proofs mechanically checked by Coq

**FSCQ:** the first certified crash-safe file system

- Basic Unix-like file system (not parallel)
- Simple specification for a subset of POSIX (e.g., no `fsync`)
- About 1.5 years of work, including learning Coq



FSCQ runs standard Unix programs: `mv`, `git`, `make`, ...

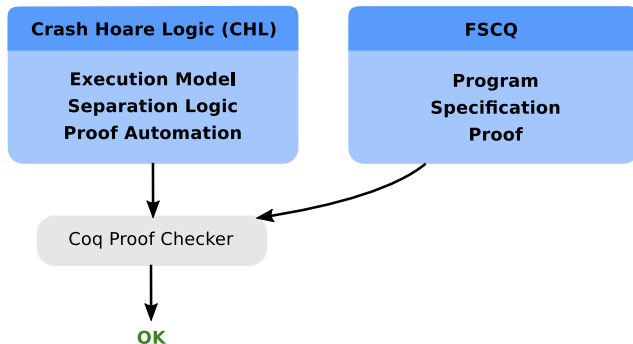
**Crash Hoare Logic (CHL)**

**Execution Model  
Separation Logic  
Proof Automation**

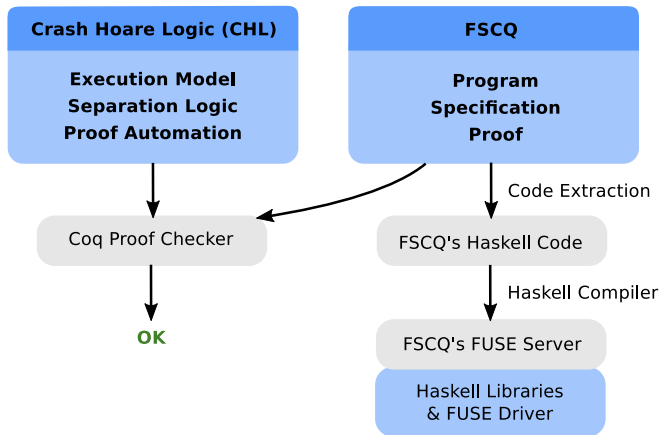
**FSCQ**

**Program  
Specification  
Proof**

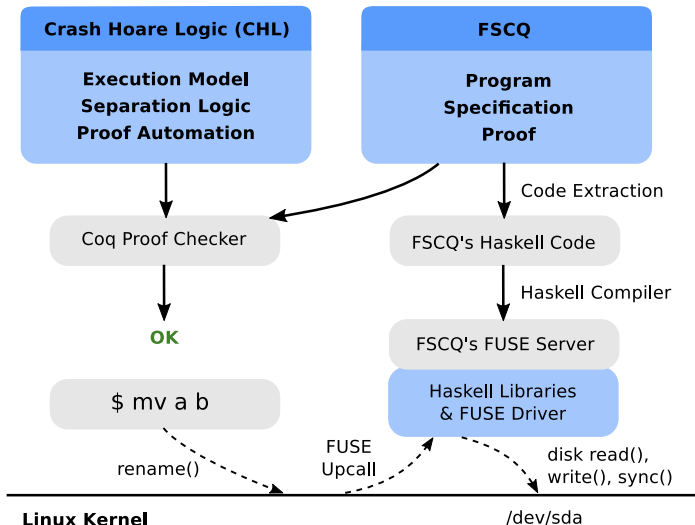
# FSCQ runs standard Unix programs: `mv`, `git`, `make`, ...



# FSCQ runs standard Unix programs: `mv`, `git`, `make`, ...



# FSCQ runs standard Unix programs: mv, git, make, ...



# How to specify what is “correct”?

Need a specification of “correct” behavior before we can prove anything

Look it up in the POSIX standard?

## How to specify what is “correct”?

Need a specification of “correct” behavior before we can prove anything

Look it up in the POSIX standard?

*[...] a power failure [...] can cause data to be lost. The data may be associated with a file that is still open, with one that has been closed, with a directory, or with any other internal system data structures associated with permanent storage. This data can be lost, in whole or part, so that only careful inspection of file contents could determine that an update did not occur.*

IEEE Std 1003.1, 2013 Edition

POSIX is vague about crash behavior

- POSIX's goal was to specify “common-denominator” behavior
- File system implementations have different interpretations
- Leads to bugs in higher-level applications [Pillai et al. OSDI'14]

# This work: “correct” is transactional

Run every file-system call inside a transaction

```
def create(d, name):  
    log_begin()  
    newfile = allocate_inode()  
    newfile.init()  
    d.add(name, newfile)  
    log_commit()
```

## This work: “correct” is transactional

Run every file-system call inside a transaction

```
def create(d, name):  
    log_begin()  
    newfile = allocate_inode()  
    newfile.init()  
    d.add(name, newfile)  
    log_commit()
```

`log_begin` and `log_commit` implement a write-ahead log on disk

After crash, replay any committed transaction in the write-ahead log



# This work: “correct” is transactional

Run every file-system call inside a transaction

```
def create(d, name):  
    log_begin()  
    newfile = allocate_inode()  
    newfile.init()  
    d.add(name, newfile)  
    log_commit()
```

`log_begin` and `log_commit` implement a write-ahead log on disk

After crash, replay any committed transaction in the write-ahead log

Q: How to formally specify both normal-case and crash behavior?

Q: How to specify that it's safe to crash during recovery itself?

# Approach: Hoare Logic specifications

$\{\text{pre}\} \text{ code } \{\text{post}\}$

**SPEC**  $\text{disk\_write}(a, v)$

**PRE**  $a \mapsto v_0$

**POST**  $a \mapsto v$

# CHL extends Hoare Logic with crash conditions

$\{\text{pre}\} \text{ code } \{\text{post}\}$   
 $\{\text{crash}\}$

<b>SPEC</b>	$\text{disk\_write}(a, v)$
<b>PRE</b>	$a \mapsto v_0$
<b>POST</b>	$a \mapsto v$
<b>CRASH</b>	$a \mapsto v_0 \vee a \mapsto v$

CHL's disk model matches what most other file systems assume:

- writing a single block is an atomic operation
- no data corruption

Disk model axiom specs: `disk_write`, `disk_read`, and `disk_sync`

# Certifying larger procedures

pre


```
def bmap(inode, bnum):  
    if bnum >= NDIRECT:  
        indirect = log_read(inode.blocks[NDIRECT])  
        return indirect[bnum - NDIRECT]  
    else:  
        return inode.blocks[bnum]
```

post

crash

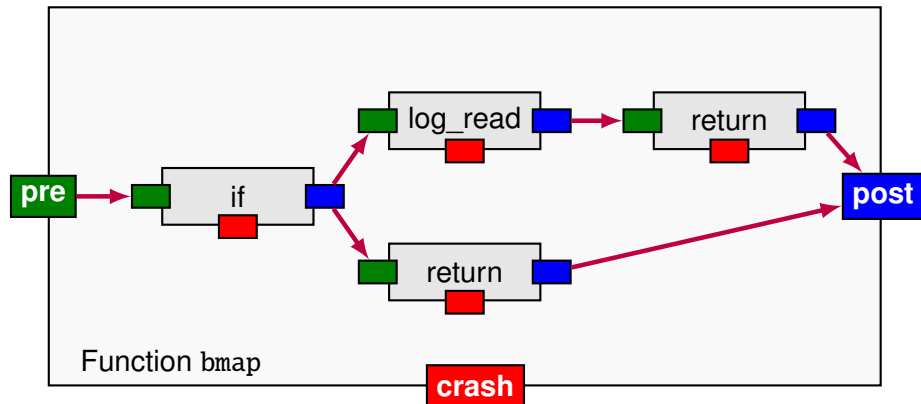
# Certifying larger procedures

Need pre/post/crash conditions for each called procedure




# Certifying larger procedures

CHL's proof automation chains pre- and postconditions



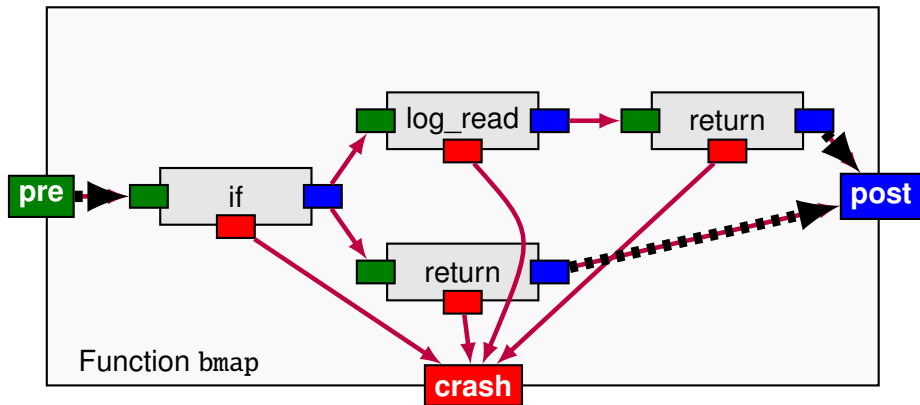
# Certifying larger procedures

CHL's proof automation combines crash conditions



# Certifying larger procedures

Remaining proof effort: changing representation invariants





## Common pattern: representation invariant

**SPEC**     $\text{log\_write}(a, v)$

**PRE**     **disk:**  $\text{log\_rep}(\text{ActiveTxn}, \text{start\_state}, \text{old\_state})$

**old\_state:**  $a \mapsto v_0$

**POST**    **disk:**  $\text{log\_rep}(\text{ActiveTxn}, \text{start\_state}, \text{new\_state})$

**new\_state:**  $a \mapsto v$

**CRASH**   **disk:**  $\text{log\_rep}(\text{ActiveTxn}, \text{start\_state}, \text{any})$

$\text{log\_rep}$  is a representation invariant

- Connects logical transaction state to an on-disk representation
- Describes the log's on-disk layout using many  $\mapsto$  primitives

# Specifying log recovery

<b>SPEC</b>	<code>log_recover()</code>
<b>PRE</b>	<b>disk:</b> <code>log_intact(<i>last_state</i>, <i>committed_state</i>)</code>
<b>POST</b>	<b>disk:</b> <code>log_rep(NoTxn, <i>last_state</i>)</code> $\vee$ <code>log_rep(NoTxn, <i>committed_state</i>)</code>
<b>CRASH</b>	<b>disk:</b> <code>log_intact(<i>last_state</i>, <i>committed_state</i>)</code>

`log_recover` is idempotent

- Crash condition implies pre condition
- $\Rightarrow$  OK to run `log_recover` *again* after a crash

# CHL summary

Key ideas: crash conditions and recovery semantics

CHL benefit: enables precise failure specifications

- Allows for automatic chaining of pre/post/crash conditions
- Reduces proof burden

CHL cost: must write crash condition for every function, loop, etc.

- Crash conditions are often simple (above logging layer)

# Breakout room discussion questions

1. Why is the crash condition in CHL important? Think about how you would specify FSCQ using normal Hoare logic, and what would be difficult.
2. What is the difference between the “ $\rightarrow$ ” symbols in these two specs?

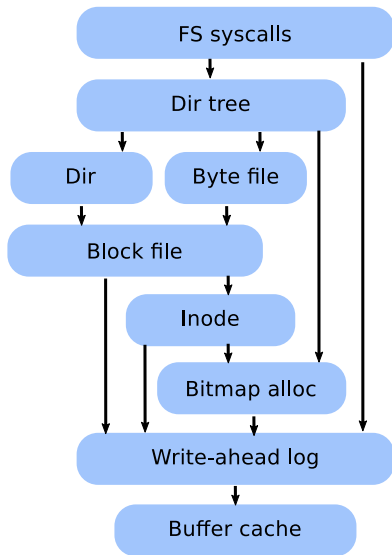
**SPEC**     `disk_write(a, v)`  
**PRE**         $a \mapsto v_0$   
**POST**         $a \mapsto v$   
**CRASH**       $a \mapsto v_0 \vee a \mapsto v$

**SPEC**     `log_write(a, v)`  
**PRE**        **disk:** `log_rep(ActiveTxn, start_state, old_state)`  
              **old\_state:**  $a \mapsto v_0$   
**POST**        **disk:** `log_rep(ActiveTxn, start_state, new_state)`  
              **new\_state:**  $a \mapsto v$   
**CRASH**      **disk:** `log_rep(ActiveTxn, start_state, any)`

# FSCQ: building a file system on top of CHL

File system design is close to v6 Unix, plus logging, minus symbolic links

Implementation aims to reduce proof effort



# Evaluation

What bugs do FSCQ's theorems eliminate?

How much development effort is required for FSCQ?

How well does FSCQ perform?

## FSCQ's theorems eliminate many bugs

One data point: once theorems proven, no implementation bugs

- Did find some mistakes in spec, as a result of end-to-end checks
- E.g., forgot to specify that extending a file should zero-fill

## FSCQ's theorems eliminate many bugs

One data point: once theorems proven, no implementation bugs

- Did find some mistakes in spec, as a result of end-to-end checks
- E.g., forgot to specify that extending a file should zero-fill

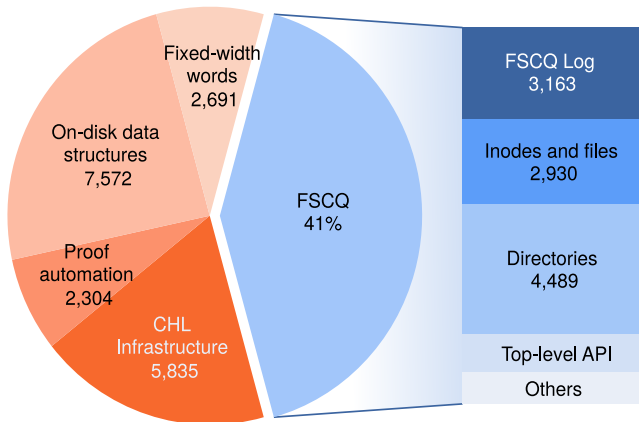
Common classes of bugs found in Linux file systems:

Bug class	Eliminated in FSCQ?
Violating file or directory invariants	Yes
Improper handling of corner cases	Yes
Returning incorrect error codes	Some
Resource-allocation bugs	Some
Mistakes in logging and recovery logic	Yes
Misusing the logging API	Yes
Bugs due to concurrent execution	No concurrency
Low-level programming errors	Yes



# Implementing CHL and FSCQ in Coq


Total of  $\sim 30,000$  lines of **verified** code, specs, and proofs  
Comparison: xv6 file system is  $\sim 3,000$  lines of code



# Change effort proportional to scope of change

- Reordering disk writes:  
~1,000 lines in FSCQLOG
- Indirect blocks:  
~1,500 lines in inode layer
- Buffer cache:  
~300 lines in FSCQLOG,  
~600 lines in rest of FSCQ
- Optimize log layout:  
~150 lines in FSCQLOG

Modest incremental effort, partially  
due to CHL's proof automation and  
FSCQ's internal layers



# Performance comparison

## File-system-intensive workload

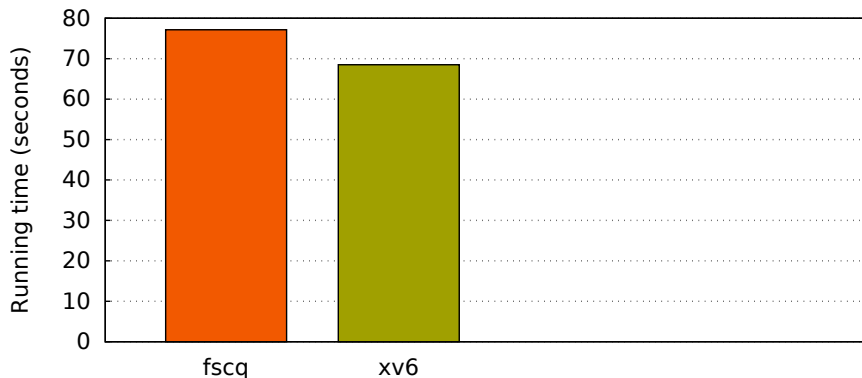
- Software development: git, make
- LFS benchmark
- mailbench: gmail-like mail server

## Compare with other (non-certified) file systems

- xv6 (similar design, written in C)
- ext4 (widely used on Linux), in non-default *synchronous* mode to match FSCQ's guarantees


## Running on an SSD on a laptop

## Running time for benchmark workload




- FSCQ slower than xv6 due to overhead of extracted Haskell

## Running time for benchmark workload



- FSCQ slower than xv6 due to overhead of extracted Haskell
- FSCQ slower than ext4 due to simple write-ahead logging design

## Opportunity: change semantics to defer durability



- FSCQ slower than xv6 due to overhead of extracted Haskell
- FSCQ slower than ext4 due to simple write-ahead logging design
- Deferred durability (ext4's default mode) allows for big improvement

## Directions for future research

Formalizing deferred durability (e.g., `fsync`)

Certifying a parallel (multi-core) file system

Certifying applications with CHL (database, key-value store, ...)

# Grading the paper: Summary

Problem: Building a provably crash-safe filesystem.

Challenges: Formal verification is difficult and time-consuming. How to specify crash behavior, such that:

- Verification is automated (as much as possible)
- Performance optimizations are possible
- Modularity is possible

Solution: Specifying crash behavior with Crash Hoare Logic (CHL).

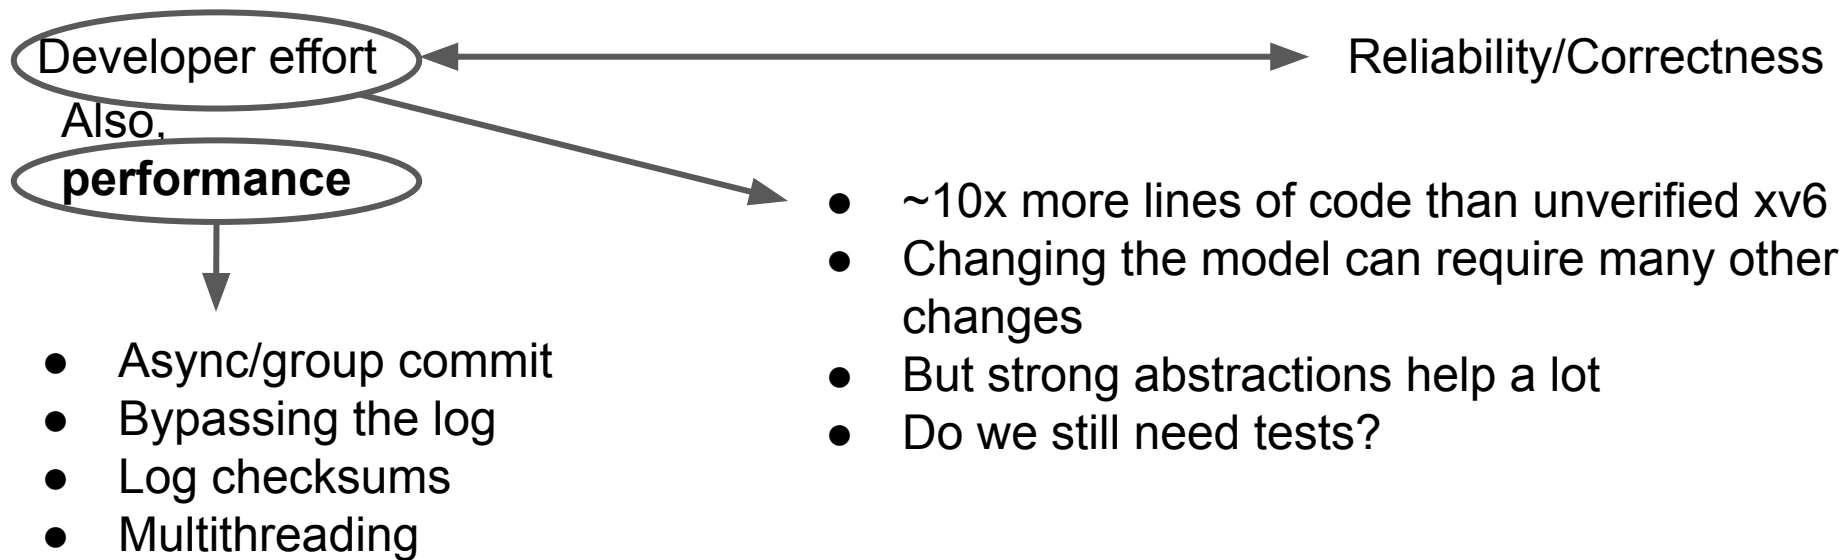


# Grading the paper: How does CHL meet the challenges?

Challenges: Formal verification is difficult and time-consuming. How to specify crash behavior, such that:

- Verification is automated (as much as possible) => automate chaining of pre/post/crash conditions; prove idempotence of recovery procedure to automate verification of e2e procedures
- Performance optimizations are possible => modeling asynchronous disk writes
- Modularity is possible => logical address spaces to abstract lower layers (e.g., FSCQLog)

# Grading the paper: Tradeoffs



In some cases, optimizations also affect higher-level filesystem spec  
=> When do the applications on top need to be verified too?